



INTRODUCTION TO HIGH PERFORMANCE COMPUTING

PERFORMANCE ANALYSIS BASICS

Radim Vavřík




EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education



MINISTRY OF EDUCATION,
YOUTH AND SPORTS

OUTLINE



Performance analysis and optimisation

- Motivation
- Hardware aspects
- Development process
- Best-practices

Performance tools and methodology


- Performance metrics
- CPU/GPU tools
- Live examples

POP CoE




Cray-1 supercomputer, source: wikipedia.org

TECHNICAL NOTES



- All presented tools/examples can be accessed and reproduced at IT4I clusters **anytime**
- Please, setup your preferred GUI access:
 1. **VNC** - server on a Karolina login node + client on laptop
 - How to? <https://docs.it4i.cz/general/accessing-the-clusters/graphical-user-interface/vnc/>
 - Recommended client <https://www.realvnc.com/en/connect/download/viewer/>
 2. **OOD** - Open OnDemand GUI via web browser, **IT4I VPN required**
 - How to? <https://docs.it4i.cz/general/accessing-the-clusters/graphical-user-interface/ood/>
 - Connection link <https://ood-karolina.it4i.cz/>
 3. **X11** - Log in via terminal with X-Window system enabled
 - How to? <https://docs.it4i.cz/general/accessing-the-clusters/graphical-user-interface/x-window-system/>
 - Usually worse UX for GUI apps due to network latency
- Most of the presented tools provide a **remote profiling**, e.g., generate output remotely from CLI while analysis can be done locally in GUI - not covered today

PERFORMANCE ANALYSIS




Who has any experience with a performance analysis tool?

- What was the tool?

Objectives today?

- Not to become an expert analyst
- Not to reach an incredible performance improvement of example codes
- Rather to get idea about the domain and introduce some tools

EFFICIENT USE OF HPC




What does it mean?

- To get the most performance out of your hardware
- The process is called **Performance Optimisation**

Why should I care about performance?

- Industry – achieve goals faster and **cheaper**
- Academia – do **more science**
 - The trend in grant competition (resource allocation) is to prove performance, scalability, etc.

KEY INGREDIENTS



Know your application

- What does it compute? (domain, methods, algorithms)
- How is it parallelized? (programming models)
- What final performance is expected? (HW limits)

Know your hardware

- What are the target machines and how many? (laptop, workstation, cluster)
- Machine-specific optimisations?


Know your tools

- Strengths and weaknesses of each tool? (easy-to-use vs detailed information)
- Learn how to use them (examples with problems/patterns)

Know your process

- Constant learning

HARDWARE ASPECTS OF PERFORMANCE



Filesystem

- I/O operations

Network

- internode communication

Memory subsystem

- NUMA effect

CPU cores

- thread/process affinity, pinning, caches


Vector registers

- vectorization, vector instructions

Accelerators

- GPU/MIC utilization, host-device data transfers

GET READY




Connect to login node via GUI

- VNC / OOD / X11

Submit an interactive job

```
| salloc --account=DD-23-116 --reservation=dd-23-116_2024-06-  
| 05T09:00:00_2024-06-05T12:30:00_5_qgpu --gpus 1
```

BASIC TOOLS



Useful to get familiar with the machine

```
| lscpu
```

```
| cat /proc/cpuinfo
```

- processor : 71 -> 72 logical processors per node
- cpu cores : 18 -> 18 physical cores per socket
- siblings : 36 -> 36 logical processors per socket
- -> 2 hyperthreads per core
- -> 2 sockets per node


```
| cat /proc/meminfo
```

- MemTotal: 196510848 kB -> 187 GB

```
| ml impi
```

```
| cpuinfo # Intel MPI utility
```

BASIC TOOLS



Use HTOP tool for interactive jobs

```
| htop -d 5 # delay 0.5s
```


- Configurable (e.g. core id, threads, process tree)

```
1 [|||||] 100.0% 10 [|||||] 100.0% 19 [|||||] 100.0% 28 [|||||] 100.0%
2 [|||||] 100.0% 11 [|||||] 100.0% 20 [|||||] 100.0% 29 [|||||] 100.0%
3 [|||||] 100.0% 12 [|||||] 100.0% 21 [|||||] 100.0% 30 [|||||] 100.0%
4 [|||||] 100.0% 13 [|||||] 100.0% 22 [|||||] 100.0% 31 [|||||] 100.0%
5 [|||||] 100.0% 14 [|||||] 100.0% 23 [|||||] 100.0% 32 [|||||] 100.0%
6 [|||||] 100.0% 15 [|||||] 100.0% 24 [|||||] 100.0% 33 [|||||] 100.0%
7 [|||||] 100.0% 16 [|||||] 100.0% 25 [|||||] 100.0% 34 [|||||] 100.0%
8 [|||||] 100.0% 17 [|||||] 100.0% 26 [|||||] 100.0% 35 [|||||] 100.0%
9 [|||||] 100.0% 18 [|||||] 100.0% 27 [|||||] 100.0% 36 [|||||] 100.0%
Mem[|||||] 13.8G/187G Tasks: 79, 346 thr; 36 running
Swp[          0K/0K] Load average: 23.62 6.93 3.32
Uptime: 15 days, 12:06:32
```

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
11171	vav0038	35	15	296M	90236	7232	R	99.5	0.0	1:02.72	./examples/wave_c 100
11203	vav0038	35	15	298M	90240	7244	R	99.5	0.0	1:03.07	./examples/wave_c 100
11212	vav0038	35	15	322M	92280	7324	R	99.5	0.0	1:03.04	./examples/wave_c 100
11162	vav0038	35	15	300M	90220	7272	R	99.5	0.0	1:03.10	./examples/wave_c 100
11188	vav0038	35	15	323M	90236	7328	R	99.5	0.0	1:03.05	./examples/wave_c 100
11207	vav0038	35	15	311M	92272	7296	R	99.5	0.0	1:03.04	./examples/wave_c 100
11164	vav0038	35	15	326M	90232	7340	R	99.5	0.0	1:03.09	./examples/wave_c 100
11195	vav0038	35	15	298M	90232	7232	R	99.5	0.0	1:03.09	./examples/wave_c 100
11158	vav0038	35	15	319M	92284	7304	R	99.5	0.0	1:03.07	./examples/wave_c 100

F1Help F2Setup F3Search F4Filter F5Tree F6SortBy F7Nice -F8Nice +F9Kill F10Quit

BASIC TOOLS




Similar tool for NVIDIA GPUs

```
| watch -n 1 nvidia-smi
```


NVIDIA-SMI 550.54.15			Driver Version: 550.54.15		CUDA Version: 12.4			
GPU	Name	Persistence-M	Bus-Id	Disp.A	Memory-Usage	Volatile	Uncorr.	ECC
	Fan	Temp				GPU-Util	Compute M.	MIG M.
0	NVIDIA A100-SXM4-40GB	N/A	Off	00000000:0B:00.0	0MiB / 40960MiB	0%	Default	0
	31C	P0	52W / 400W				Disabled	
Processes:								
GPU	GI	CI	PID	Type	Process name	GPU Memory Usage		
ID	ID							
No running processes found								

PERFORMANCE-AWARE DEVELOPMENT PROCESS




1. Develop correct functionality (testing helps)
2. Identify bottlenecks (performance limiters) using performance tools
3. Optimise bottlenecks until satisfied
 1. Build a hypothesis (ask a question)
 2. Explain the behavior (answer the question)
 3. Change the code (double-check correct functionality)
 4. Verify optimisations using profiling tools
4. Repeat until job done

BEST PRACTICES




- Do not optimise your code prematurely!
- Focus on main computational time-consuming phases (hotspots), omit preprocessing/postprocessing phases if applicable
- The 80/20 rule:
 - Programs typically spend 80% of their time in 20% of the code
 - Programmers typically spend 20% of their effort to get 80% of the total speedup possible for the application
- Keep track of your optimisation progress over time
- Always use compute nodes for profiling (**not login nodes - shared**)
- **Use SW libraries!**

SOFTWARE LIBRARIES



General-purpose math libraries

- BLAS (MKL, OpenBLAS, ATLAS, cuBLAS, ...)
- LAPACK (MKL, OpenBLAS, ATLAS, cuSolver, ...)
- FFT (MKL, cuFFT, ...)
- ...



Domain-specific libraries


- Chemistry, Bio, Geo, Physics, CAE, Big data, ML/DL

HW-specific libraries


- GPU/MIC, Intel/AMD/IBM

Optimized implementation

- Usually much better performance than a custom code
- Do NOT reinvent a wheel!
- (But avoid overkill)



PERFORMANCE METRICS



Execution time (time, time.h, ...)

- real 0m10.245s (elapsed real time)
- user 0m19.890s (user CPU time using OMP_NUM_THREADS=2)
- sys 0m0.285s (system CPU time)

Processor speed (flop/s) and Memory throughput (GB/s)

- Calculated operations per time (e.g. $c = a + b + c \rightarrow 2$ operations)
- Transferred bytes per time (e.g. $c = a + b + c \rightarrow 3$ RD + 1 WR * 8 bytes)

Speedup and Efficiency


- $S_P = T_1 / T_P$
- $E_P = S_P / P$

Scalability

- Strong/weak scaling

Others: portability, programming ability, etc.

PEAK PERFORMANCE EXAMPLE




- The theoretical HW limits, e.g. AMD EPYC 7H12 (Rome)

Processor speed:

▪ Number of compute nodes (Karolina-size machine)	720
▪ Number of sockets (CPUs) per node	2
▪ Frequency	2.6 GHz
▪ Number of cores per socket	64
▪ FMA instructions ($\mathbf{a} * \mathbf{b} + \mathbf{c}$)	2
▪ FMA units per core	2
▪ SIMD (AVX2 256b) = 4x double precision	4
	3 833 856 Gflop/s
	3.8 Pflop/s
	(2.6 Tflop/s per socket)

PEAK PERFORMANCE EXAMPLE




Memory bandwidth:


- Number of compute nodes (Karolina-size machine)
- Number of sockets (CPUs) per node
- # channels per socket
- DDR4 bus width
- Frequency

720
2
8
8 B
3200 MT/s


294 912 000 MB/s
294 TB/s
(204 GB/s per socket)




ROOFLINE MODEL



- Shows the performance of an algorithm (application) with respect to the HW limits of the architecture
- Identify if an algorithm is **compute bound** or **memory bound**
- Based on **Operational intensity** - a ratio of FLOPS (arithmetic operations) performed with required amount of data (operands)



SPEEDUP



- **Speedup** – a ratio of a serial execution time to a parallel execution time


speedup on N processors -

$$S_N = \frac{T_1}{T_N}$$


- execution time on 1 processor

- execution time on N processors

- Linear speedup $S_N = N$
- Sub-linear speedup $S_N < N$
 - Communication
 - Load imbalance
 - Decomposition overhead
- Super-linear speedup $S_N > N$
 - Cache
 - Algorithm




SCALABILITY




- **Scalability** - the ability to maintain performance gain when system and problem size increase
- **Amdahl's law** – maximum achievable speedup is limited by the serial portion of the code





$$S_{MAX} = \frac{1}{\frac{Y}{100}}$$




SCALABILITY





- **Strong scaling** - how the processing time varies with the number of processors for a **fixed total problem size**




SCALABILITY



- **Weak scaling** - how the processing time varies with the number of processors for a **fixed problem size per processing unit**



CLASSIFICATION OF PERFORMANCE TOOLS



- There are many tools that can be classified by the implemented approach

Data collecting mechanism


- **Sampling** - automatically collect data per time unit
- **Instrumentation** - manually/automatically add instructions to the source code to collect data - intrusive

Form of data presentation

- **Reports** - general overview of the whole application
- **Profiling** - accumulated characteristics of metrics
- **Tracing** - details about selected events - intrusive

Analysis of the collected data


- **Online** - during the execution - rare
- **Post mortem** - after the execution



Example of a trace, source: tools.bsc.es

Modeling - simulate state, ideal network, HW failure, etc.

PERFORMANCE TOOLS - CPU



- Single-node/parallel, architecture, language, programming model, focus (instrumentation, correctness checking, etc.)


Proprietary tools – licenses usually available on clusters

- Linaro (~~ARM (Allinea)~~) Performance Report
- Linaro (~~ARM (Allinea)~~) MAP
- Intel Application Performance Snapshot
- Intel Vtune
- AMD µProf
- Vampir

Open-source tools (VI-HPS)

- BSC tools (Extrae/Paraver)
- JSC tools (Score-P/Scalasca/Cube)
- MAQAO
- <https://www.vi-hps.org/tools/tools.html> (guide)

GPU PROFILING – TOOLS




GUI tools

- NVIDIA Nsight Systems – **system-level** profiling
- NVIDIA Nsight Compute – **CUDA kernel-level** profiling
- NVIDIA Visual Profiler - deprecated

Command-line tools - for no GUI access


(e.g. in batch jobs)

- NVIDIA nsys
- NVIDIA ncu
- AMD ROC-profiler
 - analogous to nsys
 - Chrome for visualization
- NVIDIA nvprof
 - deprecated



Nsight tools, source: nvidia.com

ARM PERFORMANCE REPORTS



Compute

MPI


I/O

- Global high-level overview of the application
- No source code or recompilation required
- Run: **perf-report srun -n <#procs> <app>**
- Auto-generated text and HTML output
- Report summary (Compute, MPI, Input/Output)
- CPU, MPI, I/O, OpenMP, Memory, Energy, Accelerator breakdown sections
- Advanced configuration through command line flags possible

arm
PERFORMANCE
REPORTS

Command: mpirun -np 8 examples/wave_openmp 60
Resources: 1 node (8 physical, 8 logical cores per node)
Memory: 15 GiB per node
Tasks: 8 processes, OMP_NUM_THREADS was 2
Machine: mars
Start time: Tue Nov 7 2017 15:35:50 (UTC)
Total time: 61 seconds (about 1 minutes)
Full path: /scratch/user/reports/examples

Summary: wave_openmp is **Compute-bound** in this configuration



Time spent running application code. High values are usually good.
This is **high**; check the CPU performance section for advice.
Time spent in MPI calls. High values are usually bad.
This is **low**; this code may benefit from a higher process count.
Time spent in filesystem I/O. High values are usually bad.
This is **negligible**; there's no need to investigate I/O performance.

This application run was **Compute-bound**. A breakdown of this time and advice for investigating further is in the **CPU** section below.

As little time is spent in **MPI** calls, this code may also benefit from running at larger scales.

CPU

A breakdown of the **72.6%** CPU time:

Category	Time Spent (%)
Single-core code	8.2%
OpenMP regions	91.8%
Scalar numeric ops	5.1%
Vector numeric ops	0.0%
Memory accesses	56.9%

The per-core performance is memory-bound. Use a profiler to identify time-consuming loops and check their cache performance.

No time is spent in **vectorized instructions**. Check the compiler's vectorization advice to see why key loops could not be vectorized.

MPI

A breakdown of the **27.4%** MPI time:

Category	Time Spent (%)
Time in collective calls	1.2%
Time in point-to-point calls	98.8%
Effective process collective rate	19.5 kB/s
Effective process point-to-point	305 kB/s

Most of the time is spent in **point-to-point calls** with a very low transfer rate. This suggests load imbalance is causing synchronization overhead; use an MPI profiler to investigate.

I/O

A breakdown of the **0.0%** I/O time:

Category	Time Spent (%)
Time in reads	0.0%
Time in writes	0.0%
Effective process read rate	0.00 bytes/s
Effective process write rate	0.00 bytes/s

No time is spent in **I/O** operations. There's nothing to optimize here!

OpenMP

A breakdown of the **91.8%** time in OpenMP regions:

Category	Time Spent (%)
Computation	9.9%
Synchronization	90.1%
Physical core utilization	100.0%
System load	167.0%

Significant time is spent **synchronizing** threads in parallel regions. Check the affected regions with a profiler.

The system load is high. Ensure background system processes are not running.

Memory

Per-process memory usage may also affect scaling:

Mean process memory usage	38.6 MiB
Peak process memory usage	53.7 MiB
Peak node memory usage	17.0%

The **peak node memory usage** is very low. Running with fewer MPI processes and more data on each process may be more efficient.


Energy

A breakdown of how energy was used:

CPU	not supported
System	not supported
Mean node power	not supported
Peak node power	0.00 W

Energy metrics are not available on this system.
CPU metrics are not supported (no intel_rapl module)

ARM PERFORMANCE REPORTS - EXAMPLE




```
| ml Forge/23.1.2 OpenMPI/4.1.6-GCC-12.2.0-CUDA-12.4.0
| ml show Forge
| cp -r /apps/all/Forge/23.1.2/examples ~/forge_examples
| cd ~/forge_examples
| make

| srun -n 16 ./wave_c 10

| mkdir perf_reports && cd perf_reports
| perf-report srun -n 16 ../wave_c 10
| firefox wave_c_16p_1n_YYYY-MM-DD_hh-mm.html &           # on login node
| OMP_NUM_THREADS=8 perf-report srun -n 2 -c 8 ../wave_openmp 10
| firefox wave_openmp_2p_1n_8t_YYYY-MM-DD_hh-mm.html &
```

ARM MAP



- Low overhead sampling profiler for localisation of bottlenecks
 - No recompilation required, only debugging symbols are useful (-g)
1. Metrics view (CPU, MPI, I/O, memory, vectorization)
 2. Source code viewer
 3. Selected lines view
 4. Output, files, callpaths
 5. Sparkline charts


| map

| map srun -n <#procs> <app> [args]


| map --profile srun -n <#procs> ...

| map <profile.map>


| perf-report <profile.map>




ARM MAP



- All charts are timelines
 - Horizontal axis time
- Vertical axis are processes
- Useful code is green
- MPI is blue
- breakout recalculated when zooming
- Multiple presets available
 - CPU
 - MPI
 - I/O
 - memory
 - ...




ARM MAP - EXAMPLE




```
| ml Forge/23.1.2 OpenMPI/4.1.6-GCC-  
| 12.2.0-CUDA-12.4.0  
  
| mkdir ~/forge_examples/map && cd  
| ~/forge_examples/map  
  
| OMP_NUM_THREADS=8 map srun -n 2 -c 8  
| ./wave_openmp 10
```

- Optionally limit duration
- Optionally adapt metrics
- Click Run
- Use the User guide!




ARM MAP - EXAMPLE




- A large section of blue means all the processes in MPI calls - try to reduce these. Triangular shape indicates load imbalance.




- A large section of dark green means all the processes in single-threaded computations - try to avoid.



- A large sections of light green - OpenMP regions being effectively used across all processes simultaneously.




NVIDIA NSIGHT SYSTEMS



Scalable system-wide performance analysis tool


- Low-overhead multi-node, multi-GPU profiling
- Visualize millions of events on a very fast GUI timeline
- Assess on timeline to narrow down frames/areas of the app to focus
- Locate optimization opportunities, CPU/GPU bottlenecks
 - or gaps of unused CPU and GPU time - idle
- Balance your workload across multiple CPUs and GPUs
- Expert system GPU utilization analysis
- Detailed information, documentation, free download
<https://developer.nvidia.com/nsight-systems>

NVIDIA NSIGHT SYSTEMS




Multi-level information


- CPU cores utilization
- MPI calls
- OpenMP, Pthreads
- OS runtime calls
- NVTX
- CUDA API calls
- HtD / DtH data transfers
- CUDA kernels / streams
- OpenACC
- CUDA libraries (cuBLAS, ...), GPU HW metrics, UCX, NIC, ...



NVIDIA NSIGHT SYSTEMS



PROFILING WITH NSIGHT SYSTEMS




GUI profiling and analysis

- ```
| ml CUDA/12.4.0 Qt5
| nsys-ui # On the allocated GPU compute node
 - File -> New Project
 - Select target for profiling... -> acnXX.karolina.it4i.cz (allocated GPU node)
 - Enter binary (absolute path with arguments if necessary)
 - Select tracing modules (CPU, OS, CUDA, GPU, NVTX,...)
 - Start
```

## Cmd line profiling + GUI analysis

- ```
| nsys profile -t cuda,osrt --stats=true --gpu-metrics-device=0
  -o profile_name ./program
| nsys-ui # On login node
  - File -> Open -> Select profile_name.nsys-rep
```

NVIDIA NSIGHT SYSTEMS - EXAMPLE



```
| git clone https://code.it4i.cz/training/intro2hpc.git  
| ml CUDA/12.4.0  
| cd hpcintro24/intro2hpc/5_gpu_accelerators/handson  
| vim vector_add.solution.cu # int count = 123456789;  
| nvcc -g -O2 -gencode arch=compute_80,code=sm_80  
vector_add.solution.cu -o vector_add  
| ./vector_add
```

Barbora sm_70
Karolina sm_80

- Perform profiling of `vector_add` example:


```
| nsys profile -t cuda,osrt --stats=true --gpu-metrics-  
device=0 -o vector_add ./vector_add
```

- An extra CUDA examples:

```
| git clone https://github.com/NVIDIA/cuda-samples.git
```

POP COE

- An EuroHPC **Centre of Excellence** (CoE)
 - On **Performance Optimisation and Productivity**
 - Promoting **best practices in parallel programming**
- Providing **FREE Services** for EU **academic AND industrial codes in all domains!**
 - **Performance Assessment:** initial analysis to identify performance issues and recommend approaches to address them
 - **Proof-of-concept:** explore the potential benefit of proposed optimisations by applying them to selected regions of the applications
 - **Correctness-check:** evaluate the correctness of hybrid MPI + OpenMP applications
 - **Energy-efficiency study:** investigate improvements of energy consumption or efficiency
 - **Advisory study:** ongoing consultancy for customers that choose to implement proposed optimisations on their own



Barcelona
Supercomputing
Center

Centro Nacional de Supercomputación



JÜLICH
Forschungszentrum

JÜLICH
SUPERCOMPUTING
CENTRE



www.pop-coe.eu



pop@bsc.es



[@POP_HPC](#)




youtube.com/POPHPC

VSB TECHNICAL
UNIVERSITY
OF OSTRAVA

IT4INNOVATIONS
NATIONAL SUPERCOMPUTING
CENTER

USEFUL LINKS



VI-HPS – Association of institutions developing tools and providing training

- Overview of the tools with a description: <https://www.vi-hps.org/cms/upload/material/general/ToolsGuide.pdf>

Nvidia tools for GPUs: Nsight Systems and Nsight Compute

Intel performance tools: VTune and Advisor

Database of code patterns and best practices developed in POP: [co-design](#)

Docs + further reading:

- https://docs.linaroforge.com/23.1.2/html/forge/performance_reports/index.html
- <https://docs.linaroforge.com/23.1.2/html/forge/map/index.html>
- <https://software.intel.com/content/www/us/en/develop/articles/intel-advisor-roofline.html>



Radim Vavřík
radim.vavrik@vsb.cz

IT4Innovations National Supercomputing Center
VSB – Technical University of Ostrava
Studentská 6231/1B
708 00 Ostrava-Poruba, Czech Republic
www.it4i.cz

VSB TECHNICAL
UNIVERSITY
OF OSTRAVA


IT4INNOVATIONS
NATIONAL SUPERCOMPUTING
CENTER




EUROPEAN UNION
European Structural and Investment Funds
Operational Programme Research,
Development and Education




ANALYSIS WITH NSIGHT SYSTEMS




Only small portion of application accelerated (for real-world apps)




ANALYSIS WITH NSIGHT SYSTEMS




GPU idle/low utilization of detailed zoom (because of Pthread creation)




ANALYSIS WITH NSIGHT SYSTEMS



Fusion opportunities: CPU launch cost + small GPU work size -> GPU idle




ANALYSIS WITH NSIGHT SYSTEMS



cudaMemcpyAsync behaving synchronously – DtH pageable memory ->
Mitigate with pinned memory



ANALYSIS WITH NSIGHT SYSTEMS



GPU idle caused by stream synchronization

